Dear Editor and Referees,

I really appreciate just how closely you have read the manuscript.

I am yet again grateful for your very constructive feedback on this manuscript. I have ...

What hasn't changed ...

There is one significant change that is not noted below ...

I have once again written a consolidated response. I have grouped your responses together where they are related and responded just once wherever possible, with my responses in italics.

With sincere gratitude,
Kristy

1 Clarify "non-unitary" interpretation of the model

Editor Point 1:
The "non-unitary" interpretation of your model is not entirely clear to me. When you say that lobbying can change the identity of the median legislator, what mechanism do you have in mind? Does lobbying influence the composition of the legislature, by affecting who gets elected? Or does lobbying somehow change the identity of the median legislator without affecting the composition of the legislature, i.e. without affecting elections? Your current wording suggests the latter, but I am not quite sure. And if it is the latter that you have in mind, how can lobbying change the identity of the median legislator without affecting the distribution of legislators? Perhaps by affecting the decision making process, e.g. who gets to sit on the relevant committees? And if so, is it appropriate to refer to the decisive legislator as the "median" legislator? All of this needs to be clarified.

- *I have clarified the interpretation after Expression (2) in Section 2.2. I do imagine the latter interpretation, but I think in some situations the former makes sense. I have a project currently underway in which I take the electoral approach in explaining the gradual reduction in trade barriers over the post-WWII period. Here, I think it is more of a stretch because we would have to imagine a new election for each enforcement phase of the game.*

2 Clarify definition of break tariff

Referee 2 Point 1:
It seems that that (7) implicitly assumes that $\tau^R(e_b) > \tau^a$. Is that right? If so, please be explicit.

- *I have added a clause to this effect at the end of the paragraph following the un-numbered equation just above Equation (7) where $\tau^R(e_b)$ is introduced.*

3 Clarify statement of lobby's incentive constraint

Referee 2 Point 3:
I believe (8) only need hold for $e_b \geq \overline{e}(\tau^a)$. I believe (8) would fail for $e_b = \overline{e}(\tau^a) - \varepsilon$, right? This seems confusing because (8) seems to be stated for any e_b.

Referee 2 Point 5:
I found the program on page 18 a little confusing. Equation (11) puts a constraint on e_b, essentially requiring that $e_b \geq \overline{e}(\tau^a)$ if I understand. Equation (10) seems to be directed toward a value for e_b such as $e_a(\tau^a)$; that is used on the equilibrium path. If we were to take the program as a mathematical object, (11) would seem to define the range of e_b that is to be considered in (10). I don't think that is what is intended. I think it might be easier to define the program using $\overline{e}(\tau^a)$ instead of waiting to define that function later.

Editor Point 2:
I think the way you write and analyse the lobby's incentive constraint is still not very clear. I would explain things in the following sequence: (a) the relevant level of e_b in constraint (8) is the best deviation effort, i.e. the optimal effort conditional on inducing a break of the agreement; (b) the best deviation effort is given by $\max\{e_{tw}, \overline{e}(\boldsymbol{\tau^a})\}$; (c) in order for a non-trivial agreement to be possible, there must exist some $\tau^a < \tau^{tw}$ such that $e_{tw} < \overline{e}(\boldsymbol{\tau^a})$; (d) to avoid any confusion, I would write constraint (13) evaluated at $\max\{e_{tw}, \overline{e}(\boldsymbol{\tau^a})\}$, rather than at $\overline{e}(\tau^a)$. Part of the reason I found the current exposition confusing is that your constraint (13) assumes $e_{tw} < \overline{e}(\tau^a)$, but this restriction is imposed and discussed only after writing constraint (13).

- *I have done exactly as you suggest just following constraint (8) and I have written constraint (11) evaluated at* $\max\{e_{tw}, \overline{e}(\boldsymbol{\tau^a})\}$ *to address Referee 2's suggestions as well. I did not write equation (13) evaluated at this* $\max\{e_{tw}, \overline{e}(\boldsymbol{\tau^a})\}$ *because it is meant to describe how the executives choose $\boldsymbol{\tau^a}$, which necessitates raising $\boldsymbol{\tau^a}$ high enough to ensure that $\overline{e}(\boldsymbol{\tau^a}) > e_{tw}$. I instead imposed and discussed the restriction in the paragraph preceding the one that describes expression (13).*

Referee 2 Point 6:
Equations (12) and (13) finally define the IC constraints using $\overline{e}(\tau^a)$. In comparison to (7) - (10), (12) replaces $\tau^R(e_b)$ with $\tau^b(\overline{e})$. Is there a difference between the τ^R and τ^b functions? If these are referring to the same functions, I suggest sticking with the former notation so as to minimize the burden on the reader. A similar remark applies to (13) in relation to (8) - (11).

- *Done throughout, including appendices.*

4 Better explain argument involving lobby's choice of effort level

Editor Point 3:
Probably I am missing something, but I find the argument in paragraph 4 of page 19 confusing. You say that in the punishment phase the lobby chooses some effort level weakly higher than e_{tw}. But consider an extremely high level of effort: how can this be credible? And how is this statement consistent with (7) and (8), where the punishment effort is e_{tw}? And don't you need to pin down the exact level of effort in the punishment phase in order to write the incentive constraints?

- *I've changed the passage to specify that the lobby chooses exactly e_{tw} and that the legislature best responds. In an earlier version with more complicated punishments, the more general specification helped to simplify the 'punishments for the punishments' layer of analysis, but this is no longer necessary (equilibrium behavior was still the same, but it's quite true that I did not state this clearly). Thank you for pointing this out.*

5 Revisit discussion of strong versus weak bindings

Editor Point 4:
The discussion of weak bindings versus strong bindings in page 24 confuses me. First, the model assumes that executives have zero political economy weights, so discussing the role of the executives' political economy weights at this juncture is confusing. But more to the point, I feel that you are not being true to your model here. The way I understand it, your model predicts that trade negotiators (executives) should prefer strong bindings to weak bindings. I know this is not what we observe in reality, but you need to be upfront about the implications of the model.

- *I have rearranged the key sentence to make clear that the political economy weights do not refer to the executives:*

> The executives, as well as legislators with small political economy weights, prefer a strong-binding agreement.

I have also stated clearly that this does not appear to match reality and have moved up from the appendix the possible justification I provided there.

6 Address minimum feasible punishment length

Editor Point 5:
The discussion right after expression (14) is puzzling. If (14) is negative for all T, doesn't this mean that no tariff below τ^{tw} can be enforced? And if this is correct, why not make this point clearly - which might actually be an interesting theoretical point - rather than invoking some ad-hoc constraint on the minimum feasible punishment length, which I don't see any justification for?

- *Expression (14) gives the derivative of the lobby's incentive constraint with respect to T. That is, it says how the slackness of the incentive constraint changes as we vary T. When (14) is negative, reducing T makes the constraint more slack, or easier to satisfy. For there to be no tariff below τ^{tw} that is enforceable, we need that the lobby's constraint itself is negative. So this does not get around the—admittedly unsatisfying—need for some ad-hock constraint on the minimum feasible punishment length.*

7 Simplify statement of Result 2

Editor Point 6:
I find the statement of Result 2 unnecessarily convoluted. Why not simply state that, if the legislature and the lobby are patient enough, the optimal punishment lasts a finite number of periods?

- *Done.*

8 Make use of bold characters consistent

Editor Point 7:
Mostly you use bold characters to denote vectors, but in (9) you use a bold W to denote the sum of the home and foreign executive welfare levels. This should be avoided.

- *There were in fact two places other than for vectors of tariffs that I used bold notation. One was for the vector of discount factors and the other was the bold W. Each was only used once after the initial definition so I removed the sum / vector definitions altogether to ease the notational burden on the reader.*

9 Remove unnecessary commentary

Referee 1 Point 1:
The main drawback of the paper, however, is that it is not an easy read and the analysis is not presented well. For example, while presenting the formal model, the author inserts various informal discussions that are very lengthy and unnecessary. Most of these discussions are provided to justify the real-world relevance of the model's assumption. I think the author could greatly improve the presentation of the model by treating it as a purely theoretical model. For example, I don't think that the following paragraph from page 25 adds any insights to the theoretical discussion in section 5:

"A change in δ_L might reflect a change in firms' planning horizons, or even their operational horizonsalthough it is not entirely clear in which direction this might work for firms who are facing extinction without sufficient protection. The lobby's patience level might also change with a change in the administrative leadership of the lobby, or as a reduced form for changes in risk aversion in a model with political uncertaintya more risk-averse lobby would effectively weigh the future, uncertain gains less relative to the current, certain cost."

There are various paragraphs similar to this one that can be simply removed from the paper.

- *I deleted the paragraph that was suggested and others like it in Section 5. I also cut and condensed throughout the text. There is more that could be cut, but I kept a bit to justify the assumptions and anything that—to the best of my memory and records—was added as a result of the editorial process at the JIE.*

- *Addressing the other suggestions from this round added about one page of length. Through this tightening process, I reduced the total length by a little over four pages.*

10 Proofread thoroughly

Referee 1 Point 2:
The paper still has various typos and a thorough proofreading is needed.

- *I first proofread the manuscript closely myself and then hired a professional to go over it again.*

11 Reconsider labeling of break effort level

Referee 2 Point 2:
It is a little awkward to use e_b to call to mind a "break" when under (7) we have that e_b is set so that the ML does not break the agreement.

- *I agree that this is awkward. I have tried to come up with some terminology and notation that means "tries to disrupt the agreement but is not necessarily successful. I have not been at all successful. I apologize.*

12 Add graph?

Referee 2 Point 3:
I suggest that the author consider adding a figure with τ on the y-axis and e_b on the x-axis, and with an upward sloping line corresponding to $\tau^R(e)$. The author could then depict τ^a and τ^{tw} in ascending order on the y-axis, and similarly $e_a(\tau^a)$, e_{tw} and $\overline{e}(\tau^a)$ in ascending order on the x-axis. Given τ^a, we can find $e_a(\tau^a)$ off of the $\tau^R(e)$ curve, and similarly for τ^{tw} and e_{tw}. If the distance between $\overline{e}(\tau^a)$ and e_{tw} is large in comparison to that between e_{tw} and $e_a(\tau^a)$, then it can be seen that the cost to the lobby of going sufficiently above its ideal point, e_{tw}, could offset the future benefit of being eliciting $\tau^R(\overline{e}(\tau^a))$ and then τ^{tw} over the punishment phase. I may be mistaken, but I think this is the basic tradeoff that sits at the foundation of the analysis.

- *I found this suggestion appealing and tried to implement it but it quickly became complicated, in particular because the lobby's payoffs are in terms of $\pi(\tau^R)$, not τ^R directly. This added to the issue of aggregating across phases of the repeated game meant that the graph seemed to confuse more than clarify.*